Sequence to Sequence – Video to Text

Subhashini Venugopalan[1]   Marcus Rohrbach[2,4]   Jeff Donahue[2]   Raymond Mooney[1]
Trevor Darrell[2]   Kate Saenko[3]

[1]University of Texas at Austin   [2]University of California, Berkeley
[3]University of Massachusetts, Lowell   [4]International Computer Science Institute, Berkeley

arXiv:1505.00487v3 [cs.CV] 19 Oct 2015

Abstract

Real-world videos often have complex dynamics; and methods for generating open-domain video descriptions should be sensitive to temporal structure and allow both input (sequence of frames) and output (sequence of words) of variable length. To approach this problem, we propose a novel end-to-end sequence-to-sequence model to generate captions for videos. For this we exploit recurrent neural networks, specifically LSTMs, which have demonstrated state-of-the-art performance in image caption generation. Our LSTM model is trained on video-sentence pairs and learns to associate a sequence of video frames to a sequence of words in order to generate a description of the event in the video clip. Our model naturally is able to learn the temporal structure of the sequence of frames as well as the sequence model of the generated sentences, i.e. a language model. We evaluate several variants of our model that exploit different visual features on a standard set of YouTube videos and two movie description datasets (M-VAD and MPII-MD).

1. Introduction

Describing visual content with natural language text has recently received increased interest, especially describing images with a single sentence [8, 5, 16, 18, 20, 23, 29, 40]. Video description has so far seen less attention despite its important applications in human-robot interaction, video indexing, and describing movies for the blind. While image description handles a variable length output sequence of words, video description also has to handle a variable length input sequence of frames. Related approaches to video description have resolved variable length input by holistic video representations [29, 28, 11], pooling over frames [39], or sub-sampling on a fixed number of input frames [43]. In contrast, in this work we propose a sequence to sequence model which is trained end-to-end and is able to learn arbitrary temporal structure in the input sequence. Our model is sequence to sequence in a sense that it reads in frames
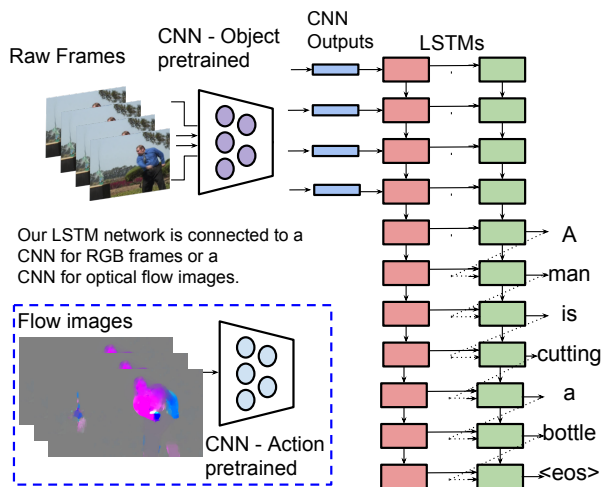

Figure 1. Our S2VT approach performs video description using a sequence to sequence model. It incorporates a stacked LSTM which first reads the sequence of frames and then generates a sequence of words. The input visual sequence to the model is comprised of RGB and/or optical flow CNN outputs.

sequentially and outputs words sequentially.

The problem of generating descriptions in open domain videos is difficult not just due to the diverse set of objects, scenes, actions, and their attributes, but also because it is hard to determine the salient content and describe the event appropriately in context. To learn what is worth describing, our model learns from video clips and paired sentences that describe the depicted events in natural language. We use Long Short Term Memory (LSTM) networks [12], a type of recurrent neural network (RNN) that has achieved great success on similar sequence-to-sequence tasks such as speech recognition [10] and machine translation [34]. Due to the inherent sequential nature of videos and language, LSTMs are well-suited for generating descriptions of events in videos.

The main contribution of this work is to propose a novel model, S2VT, which learns to directly map a sequence of

frames to a sequence of words. Figure 1 depicts our model. A stacked LSTM first encodes the frames one by one, taking as input the output of a Convolutional Neural Network (CNN) applied to each input frame's intensity values. Once all frames are read, the model generates a sentence word by word. The encoding and decoding of the frame and word representations are learned jointly from a parallel corpus. To model the temporal aspects of activities typically shown in videos, we also compute the optical flow [2] between pairs of consecutive frames. The flow images are also passed through a CNN and provided as input to the LSTM. Flow CNN models have been shown to be beneficial for activity recognition [31, 8].

To our knowledge, this is the first approach to video description that uses a general sequence to sequence model. This allows our model to (a) handle a variable number of input frames, (b) learn and use the temporal structure of the video and (c) learn a language model to generate natural, grammatical sentences. Our model is learned jointly and end-to-end, incorporating both intensity and optical flow inputs, and does not require an explicit attention model. We demonstrate that S2VT achieves state-of-the-art performance on three diverse datasets, a standard YouTube corpus (MSVD) [3] and the M-VAD [37] and MPII Movie Description [28] datasets. Our implementation (based on the *Caffe* [15] deep learning framework) is available on github. https://github.com/vsubhashini/caffe/tree/recurrent/examples/s2vt.

2. Related Work

Early work on video captioning considered tagging videos with metadata [1] and clustering captions and videos [14, 25, 42] for retrieval tasks. Several previous methods for generating sentence descriptions [11, 19, 36] used a two stage pipeline that first identifies the semantic content (subject, verb, object) and then generates a sentence based on a template. This typically involved training individual classifiers to identify candidate objects, actions and scenes. They then use a probabilistic graphical model to combine the visual confidences with a language model in order to estimate the most likely content (subject, verb, object, scene) in the video, which is then used to generate a sentence. While this simplified the problem by detaching content generation and surface realization, it requires selecting a set of relevant objects and actions to recognize. Moreover, a template-based approach to sentence generation is insufficient to model the richness of language used in human descriptions – e.g., which attributes to use and how to combine them effectively to generate a good description. In contrast, our approach avoids the separation of content identification and sentence generation by learning to directly map videos to full human-provided sentences, learning a language model simultaneously conditioned on visual features.

Our models take inspiration from the image caption generation models in [8, 40]. Their first step is to generate a fixed length vector representation of an image by extracting features from a CNN. The next step learns to decode this vector into a sequence of words composing the description of the image. While any RNN can be used in principle to decode the sequence, the resulting long-term dependencies can lead to inferior performance. To mitigate this issue, LSTM models have been exploited as sequence decoders, as they are more suited to learning long-range dependencies. In addition, since we are using variable-length video as input, we use LSTMs as sequence to sequence transducers, following the language translation models of [34].

In [39], LSTMs are used to generate video descriptions by pooling the representations of individual frames. Their technique extracts CNN features for frames in the video and then mean-pools the results to get a single feature vector representing the entire video. They then use an LSTM as a sequence decoder to generate a description based on this vector. A major shortcoming of this approach is that this representation completely ignores the ordering of the video frames and fails to exploit any temporal information. The approach in [8] also generates video descriptions using an LSTM; however, they employ a version of the two-step approach that uses CRFs to obtain semantic tuples of activity, object, tool, and locatation and then use an LSTM to translate this tuple into a sentence. Moreover, the model in [8] is applied to the limited domain of cooking videos while ours is aimed at generating descriptions for videos "in the wild".

Contemporaneous with our work, the approach in [43] also addresses the limitations of [39] in two ways. First, they employ a 3-D convnet model that incorporates spatio-temporal motion features. To obtain the features, they assume videos are of fixed volume *(width, height, time)*. They extract dense trajectory features (HoG, HoF, MBH) [41] over non-overlapping cuboids and concatenate these to form the input. The 3-D convnet is pre-trained on video datasets for action recognition. Second, they include an attention mechanism that learns to weight the frame features non-uniformly conditioned on the previous word input(s) rather than uniformly weighting features from all frames as in [39]. The 3-D convnet alone provides limited performance improvement, but in conjunction with the attention model it notably improves performance. We propose a simpler approach to using temporal information by using an LSTM to encode the sequence of video frames into a distributed vector representation that is sufficient to generate a sentential description. Therefore, our direct sequence to sequence model does not require an explicit attention mechanism.

Another recent project [33] uses LSTMs to predict the future frame sequence from an encoding of the previous frames. Their model is more similar to the language translation model in [34], which uses one LSTM to encode the

input text into a fixed representation, and another LSTM to decode it into a different language. In contrast, we employ a single LSTM that learns both encoding and decoding based on the inputs it is provided. This allows the LSTM to share weights between encoding and decoding.

Other related work includes [24, 8], which uses LSTMs for activity classification, predicting an activity class for the representation of each image/flow frame. In contrast, our model generates captions after encoding the complete sequence of optical flow images. Specifically, our final model is an ensemble of the sequence to sequence models trained on raw images and optical flow images.

3. Approach

We propose a sequence to sequence model for video description, where the input is the sequence of video frames $(x_1, \ldots, x_n)$, and the output is the sequence of words $(y_1, \ldots, y_m)$. Naturally, both the input and output are of variable, potentially different, lengths. In our case, there are typically many more frames than words.

In our model, we estimate the conditional probability of an output sequence $(y_1, \ldots, y_m)$ given an input sequence $(x_1, \ldots, x_n)$ i.e.

$$p(y_1, \ldots, y_m | x_1, \ldots, x_n) \quad (1)$$

This problem is analogous to machine translation between natural languages, where a sequence of words in the input language is translated to a sequence of words in the output language. Recently, [6, 34] have shown how to effectively attack this sequence to sequence problem with an LSTM Recurrent Neural Network (RNN). We extend this paradigm to inputs comprised of sequences of video frames, significantly simplifying prior RNN-based methods for video description. In the following, we describe our model and architecture in detail, as well as our input and output representation for video and sentences.

3.1. LSTMs for sequence modeling

The main idea to handle variable-length input and output is to first encode the input sequence of frames, one at a time, representing the video using a latent vector representation, and then decode from that representation to a sentence, one word at a time.

Let us first recall the Long Short Term Memory RNN (LSTM), originally proposed in [12]. Relying on the LSTM unit proposed in [44], for an input x_t at time step t, the LSTM computes a hidden/control state h_t and a memory cell state c_t which is an encoding of everything the cell has observed until time t:

$$
\begin{aligned}
i_t &= \sigma(W_{xi}x_t + W_{hi}h_{t-1} + b_i) \\
f_t &= \sigma(W_{xf}x_t + W_{hf}h_{t-1} + b_f) \\
o_t &= \sigma(W_{xo}x_t + W_{ho}h_{t-1} + b_o) \\
g_t &= \phi(W_{xg}x_t + W_{hg}h_{t-1} + b_g) \\
c_t &= f_t \odot c_{t-1} + i_t \odot g_t \\
h_t &= o_t \odot \phi(c_t)
\end{aligned}
\quad (2)
$$

where σ is the sigmoidal non-linearity, ϕ is the hyperbolic tangent non-linearity, $\odot$ represents the element-wise product with the gate value, and the weight matrices denoted by W_{ij} and biases b_j are the trained parameters.

Thus, in the encoding phase, given an input sequence X $(x_1, \ldots, x_n)$, the LSTM computes a sequence of hidden states $(h_1, \ldots, h_n)$. During decoding it defines a distribution over the output sequence Y $(y_1, \ldots, y_m)$ given the input sequence X as $p(Y|X)$ is

$$p(y_1, \ldots, y_m | x_1, \ldots, x_n) = \prod_{t=1}^{m} p(y_t | h_{n+t-1}, y_{t-1}) \quad (3)$$

where the distribution of $p(y_t | h_{n+t})$ is given by a softmax over all of the words in the vocabulary (see Equation 5). Note that h_{n+t} is obtained from h_{n+t-1}, y_{t-1} based on the recursion in Equation 2.

3.2. Sequence to sequence video to text

Our approach, S2VT, is depicted in Figure 2. While [6, 34] first encode the input sequence to a fixed length vector using one LSTM and then use another LSTM to map the vector to a sequence of outputs, we rely on a single LSTM for both the encoding and decoding stage. This allows parameter sharing between the encoding and decoding stage.

Our model uses a stack of two LSTMs with 1000 hidden units each. Figure 2 shows the LSTM stack unrolled over time. When two LSTMs are stacked together, as in our case, the hidden representation (h_t) from the first LSTM layer (colored red) is provided as the input (x_t) to the second LSTM (colored green). The top LSTM layer in our architecture is used to model the visual frame sequence, and the next layer is used to model the output word sequence.

Training and Inference In the first several time steps, the top LSTM layer (colored red in Figure 2) receives a sequence of frames and encodes them while the second LSTM layer receives the hidden representation (h_t) and concatenates it with null padded input words (zeros), which it then encodes. There is no loss during this stage when the LSTMs are encoding. After all the frames in the video clip are exhausted, the second LSTM layer is fed the beginning-of-sentence (<BOS>) tag, which prompts it to start decoding its current hidden representation into a sequence of words. While training in the decoding stage, the model maximizes for the log-likelihood of the predicted output sentence given the hidden representation of the visual frame sequence, and the previous words it has seen. From Equation 3 for a model with parameters θ and output sequence $Y = (y_1, \ldots, y_m)$, this is formulated as:

$$\theta^* = \underset{\theta}{\operatorname{argmax}} \sum_{t=1}^{m} \log p(y_t | h_{n+t-1}, y_{t-1}; \theta) \quad (4)$$

This log-likelihood is optimized over the entire training dataset using stochastic gradient descent. The loss is computed only when the LSTM is learning to decode. Since this
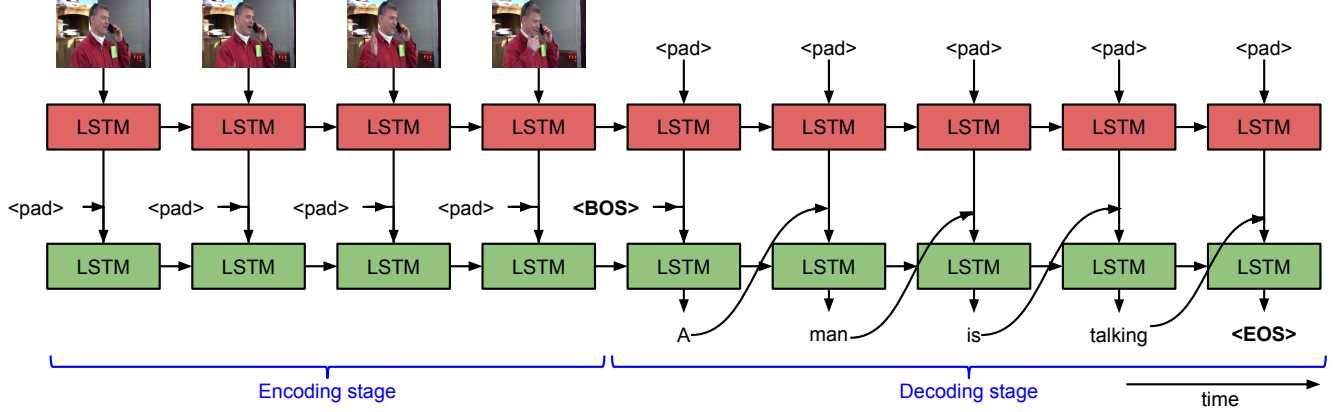
Figure 2. We propose a stack of two LSTMs that learn a representation of a sequence of frames in order to decode it into a sentence that describes the event in the video. The top LSTM layer (colored red) models visual feature inputs. The second LSTM layer (colored green) models language given the text input and the hidden representation of the video sequence. We use <BOS> to indicate begin-of-sentence and <EOS> for the end-of-sentence tag. Zeros are used as a <pad> when there is no input at the time step.

loss is propagated back in time, the LSTM learns to generate an appropriate hidden state representation (h_n) of the input sequence. The output (z_t) of the second LSTM layer is used to obtain the emitted word (y). We apply a softmax function to get the probability distribution over the words y' in the vocabulary V:

$$p(y|z_t) = \frac{\exp(W_y z_t)}{\sum_{y' \in V} \exp(W_{y'} z_t)} \quad (5)$$

We note that, during the decoding phase, the visual frame representation for the first LSTM layer is simply a vector of zeros that acts as padding input. We require an explicit end-of-sentence tag (<EOS>) to terminate each sentence since this enables the model to define a distribution over sequences of varying lengths. At test time, during each decoding step we choose the word y_t with the maximum probability after the softmax (from Equation 5) until it emits the <EOS> token.

3.3. Video and text representation

RGB frames. Similar to previous LSTM-based image captioning efforts [8, 40] and video-to-text approaches [39, 43], we apply a convolutional neural network (CNN) to input images and provide the output of the top layer as input to the LSTM unit. In this work, we report results using the output of the fc7 layer (after applying the ReLU non-linearity) on the Caffe Reference Net (a variant of AlexNet) and also the 16-layer VGG model [32]. We use CNNs that are pre-trained on the 1.2M image ILSVRC-2012 object classification subset of the ImageNet dataset [30] and made available publicly via the Caffe ModelZoo.[1] Each input video frame is scaled to 256x256, and is cropped to a random 227x227

region. It is then processed by the CNN. We remove the original last fully-connected classification layer and learn a new linear embedding of the features to a 500 dimensional space. The lower dimension features form the input (x_t) to the first LSTM layer. The weights of the embedding are learned jointly with the LSTM layers during training.

Optical Flow. In addition to CNN outputs from raw image (RGB) frames, we also incorporate optical flow measures as input sequences to our architecture. Others [24, 8] have shown that incorporating optical flow information to LSTMs improves activity classification. As many of our descriptions are activity centered, we explore this option for video description as well. We follow the approach in [8, 9] and first extract classical variational optical flow features [2]. We then create flow images (as seen in Figure 1) in a manner similar to [9], by centering x and y flow values around 128 and multiplying by a scalar such that flow values fall between 0 and 255. We also calculate the flow magnitude and add it as a third channel to the flow image. We then use a CNN [9] initialized with weights trained on the UCF101 video dataset to classify optical flow images into 101 activity classes. The fc6 layer activations of the CNN are embedded in a lower 500 dimensional space which is then given as input to the LSTM. The rest of the LSTM architecture remains unchanged for flow inputs.

In our combined model, we use a shallow fusion technique to integrate flow and RGB features. At each time step of the decoding phase, the model proposes a set of candidate words. We then rescore these hypotheses with the weighted sum of the scores by the flow and RGB networks, where we only need to recompute the score of each new word $p(y_t = y')$ as:

$$\alpha \cdot p_{rgb}(y_t = y') + (1 - \alpha) \cdot p_{flow}(y_t = y')$$

the hyper-parameter α is tuned on the validation set.

---

[1] https://github.com/BVLC/caffe/wiki/Model-Zoo

	MSVD	MPII-MD	MVAD
#-sentences	80,827	68,375	56,634
#-tokens	567,874	679,157	568,408
vocab	12,594	21,700	18,092
#-videos	1,970	68,337	46,009
avg. length	10.2s	3.9s	6.2s
#-sents per video	≈41	1	1-2

Table 1. Corpus Statistics. The the number of tokens in all datasets are comparable, however MSVD has multiple descriptions for each video while the movie corpora (MPII-MD, MVAD) have a large number of clips with a single description each. Thus, the number of video-sentence pairs in all 3 datasets are comparable.

Text input. The target output sequence of words are represented using one-hot vector encoding (1-of-N coding, where N is the size of the vocabulary). Similar to the treatment of frame features, we embed words to a lower 500 dimensional space by applying a linear transformation to the input data and learning its parameters via back propagation. The embedded word vector concatenated with the output (h_t) of the first LSTM layer forms the input to the second LSTM layer (marked green in Figure 2). When considering the output of the LSTM we apply a softmax over the complete vocabulary as in Equation 5.

4. Experimental Setup

This secction describes the evaluation of our approach. We first describe the datasets used, then the evaluation protocol, and then the details of our models.

4.1. Video description datasets

We report results on three video description corpora, namely the Microsoft Video Description corpus (MSVD) [3], the MPII Movie Description Corpus (MPII-MD) [28], and the Montreal Video Annotation Dataset (M-VAD) [37]. Together they form the largest parallel corpora with open domain video and natural language descriptions. While MSVD is based on web clips with short human-annotated sentences, MPII-MD and M-VAD contain Hollywood movie snippets with descriptions sourced from script data and audio description. Statistics of each corpus are presented in Table 1.

4.1.1 Microsoft Video Description Corpus (MSVD)

The Microsoft Video description corpus [3], is a collection of Youtube clips collected on Mechanical Turk by requesting workers to pick short clips depicting a single activity. The videos were then used to elicit single sentence descriptions from annotators. The original corpus has multi-lingual descriptions, in this work we use only the English descriptions. We do minimal pre-processing on the text by con-

verting all text to lower case, tokenizing the sentences and removing punctuation. We use the data splits provided by [39]. Additionally, in each video, we sample every tenth frame as done by [39].

4.1.2 MPII Movie Description Dataset (MPII-MD)

MPII-MD [28] contains around 68,000 video clips extracted from 94 Hollywood movies. Each clip is accompanied with a single sentence description which is sourced from movie scripts and audio description (AD) data. AD or Descriptive Video Service (DVS) is an additional audio track that is added to the movies to describe explicit visual elements in a movie for the visually impaired. Although the movie snippets are manually aligned to the descriptions, the data is very challenging due to the high diversity of visual and textual content, and the fact that most snippets have only a single reference sentence. We use the training/validation/test split provided by the authors and extract every fifth frame (videos are shorter than MSVD, averaging 94 frames).

4.1.3 Montreal Video Annotation Dataset (M-VAD)

The M-VAD movie description corpus [37] is another recent collection of about 49,000 short video clips from 92 movies. It is similar to MPII-MD, but contains only AD data with automatic alignment. We use the same setup as for MPII-MD.

4.2. Evaluation Metrics

Quantitative evaluation of the models are performed using the METEOR [7] metric which was originally proposed to evaluate machine translation results. The METEOR score is computed based on the alignment between a given hypothesis sentence and a set of candidate reference sentences. METEOR compares exact token matches, stemmed tokens, paraphrase matches, as well as semantically similar matches using WordNet synonyms. This semantic aspect of METEOR distinguishes it from others such as BLEU [26], ROUGE-L [21], or CIDEr [38]. The authors of CIDEr [38] evaluated these four measures for image description. They showed that METEOR is always better than BLEU and ROUGE and outperforms CIDEr when the number of references are small (CIDEr is comparable to METEOR when the number of references are large). Since MPII-MD and M-VAD have only a single reference, we decided to use METEOR in all our evaluations. We employ METEOR version 1.5 [2] using the code[3] released with the Microsoft COCO Evaluation Server [4].

4.3. Experimental details of our models

All our models take as input either the raw RGB frames directly feeding into the CNN, or pre-processed optical flow

---

images (described in Section 3.3). In all of our models, we unroll the LSTM to a fixed 80 time steps during training. We found this to be a good trade-off between memory consumption and the ability to provide many frames (videos) to the LSTM. This setting allows us to fit multiple videos in a single mini-batch (up to 8 for AlexNet and up to 3 for flow models). We note that 94% of the YouTube training videos satisfied this limit (with frames sampled at the rate of 1 in 10). For videos with fewer than 80 time steps (of words and frames), we pad the remaining inputs with zeros. For longer videos, we truncate the number of frames to ensure that the sum of the number of frames and words is within this limit. At test time, we do not constrain the length of the video and our model views all sampled frames. We use the pre-trained AlexNet and VGG CNNs. For VGG, we fix all layers below fc7 to reduce memory consumption and allow faster training.

We compare our sequence to sequence LSTM architecture with RGB image features extracted from both AlexNet, and the 16-layer VGG network. In order to compare features from the VGG network with previous models, we include the performance of the mean-pooled model proposed in [39] using the output of the fc7 layer from the 16 layer VGG as a baseline (line 3, Table 2). All our sequence to sequence models are referenced in Table 2 under S2VT. Our first variant, RGB (AlexNet) is the end-to-end model that uses AlexNet on RGB frames. Flow (AlexNet) refers to the model that is obtained by training on optical flow images. RGB (VGG) refers to the model with the 16-layer VGG model on RGB image frames. We also experiment with randomly re-ordered input frames (line 10) to verify that S2VT learns temporal-sequence information. Our final model is an ensemble of the RGB (VGG) and Flow (AlexNet) where the prediction at each time step is a weighted average of the prediction from the individual models.

4.4. Related approaches

We compare our sequence to sequence models against the factor graph model (FGM) in [36], the mean-pooled models in [39] and the Soft-Attention models of [43].
FGM proposed in [36] uses a two step approach to first obtain confidences on subject, verb, object and scene (S,V,O,P) elements and combines these with confidences from a language model using a factor graph to infer the most likely (S,V,O,P) tuple in the video. It then generates a sentence based on a template.
The **Mean Pool** model proposed in [39] pools AlexNet fc7 activations across all frames to create a fixed-length vector representation of the video. It then uses an LSTM to then decode the vector into a sequence of words. Further, the model ia pre-trained on the Flickr30k [13] and MSCOCO [22] image-caption datasets and fine-tuned on MSVD for a significant improvement in performance. We compare

Model	METEOR	
FGM [36]	23.9	(1)
Mean pool		
- AlexNet [39]	26.9	(2)
- VGG	27.7	(3)
- AlexNet COCO pre-trained [39]	29.1	(4)
- GoogleNet [43]	28.7	(5)
Temporal attention		
- GoogleNet [43]	29.0	(6)
- GoogleNet + 3D-CNN [43]	29.6	(7)
S2VT (ours)		
- Flow (AlexNet)	24.3	(8)
- RGB (AlexNet)	27.9	(9)
- RGB (VGG) random frame order	28.2	(10)
- RGB (VGG)	29.2	(11)
- RGB (VGG) + Flow (AlexNet)	29.8	(12)

Table 2. MSVD dataset (METEOR in %, higher is better).

our models against their basic mean-pooled model and their best model obtained from fine-tuning on Flickr30k and COCO datasets. We also compare against the GoogleNet [35] variant of the mean-pooled model reported in [43].
The **Temporal-Attention** model in [43] is a combination of weighted attention over a fixed set of video frames with input features from GoogleNet and a 3D-convnet trained on HoG, HoF and MBH features from an activity classification model.

5. Results and Discussion

This section discussses the result of our evaluation shown in Tables 2, 4, and 5.

5.1. MSVD dataset

Table 2 shows the results on the MSVD dataset. Rows 1 through 7 present related approaches and the rest are variants of our S2VT approach. Our basic S2VT AlexNet model on RGB video frames (line 9 in Table 2) achieves 27.9% METEOR and improves over the basic mean-pooled model in [39] (line 2, 26.9%) as well as the VGG mean-pooled model (line 3, 27.7%);suggesting that S2VT is a more powerful approach. When the model is trained with randomly-ordered frames (line 10 in Table 2), the score is considerably lower, clearly demonstrating that the model benefits from exploiting temporal structure.

Our S2VT model which uses flow images (line 8) achieves only 24.3% METEOR but improves the performance of our VGG model from 29.2% (line 11) to 29.8% (line 12), when combined. A reason for the low performance of the flow model could be that optical flow features even for the same activity can vary significantly with context e.g. 'panda eating' vs 'person eating'. Also, the

Edit-Distance	$k = 0$	$k <= 1$	$k <= 2$	$k <= 3$
MSVD	42.9	81.2	93.6	96.6
MPII-MD	28.8	43.5	56.4	83.0
MVAD	15.6	28.7	37.8	45.0

Table 3. Percentage of generated sentences which match a sentence of the training set with an edit (Levenshtein) distance of less than 4. All values reported in percentage (%).

model only receives very weak signals with regard to the kind of activities depicted in YouTube videos. Some commonly used verbs such as "play" are polysemous and can refer to playing a musical instrument ("playing a guitar") or playing a sport ("playing golf"). However, integrating RGB with Flow improves the quality of descriptions.

Our ensemble using both RGB and Flow performs slightly better than the best model proposed in [43], temporal attention with GoogleNet + 3D-CNN (line 7). The modest size of the improvement is likely due to the much stronger 3D-CNN features (as the difference to GoogleNet alone (line 6) suggests). Thus, the closest comparison between the Temporal Attention Model [43] and S2VT is arguably S2VT with VGG (line 12) vs. their GoogleNet-only model (line 6).

Figure 3 shows descriptions generated by our model on sample Youtube clips from MSVD. To compare the originality in generation, we compute the Levenshtein distance of the predicted sentences with those in the training set. From Table 3, for the MSVD corpus, 42.9% of the predictions are identical to some training sentence, and another 38.3% can be obtained by inserting, deleting or substituting one word from some sentence in the training corpus. We note that many of the descriptions generated are relevant.

5.2. Movie description datasets

For the more challenging MPII-MD and M-VAD datasets we use our single best model, namely S2VT trained on RGB frames and VGG. To avoid over-fitting on the movie corpora we employ drop-out which has proved to be beneficial on these datasets [27]. We found it was best to use dropout at the inputs and outputs of both LSTM layers. Further, we used ADAM [17] for optimization with a first momentum coefficient of 0.9 and a second momentum coefficient of 0.999. For MPII-MD, reported in Table 4, we improve over the SMT approach from [28] from 5.6% to 7.1% METEOR and over Mean pooling [39] by 0.4%. Our performance is similar to Visual-Labels [27], a contemporaneous LSTM-based approach which uses no temporal encoding, but more diverse visual features, namely object detectors, as well as activity and scene classifiers.

On M-VAD we achieve 6.7% METEOR which significantly outperforms the temporal attention model [43]

Approach	METEOR
SMT (best variant) [28]	5.6
Visual-Labels [27]	7.0
Mean pool (VGG)	6.7
S2VT: RGB (VGG), ours	7.1

Table 4. MPII-MD dataset (METEOR in %, higher is better).

Approach	METEOR
Visual-Labels [27]	6.3
Temporal att. (GoogleNet+3D-CNN) [43] [4]	4.3
Mean pool (VGG)	6.1
S2VT: RGB (VGG), ours	6.7

Table 5. M-VAD dataset (METEOR in %, higher is better).

(4.3%)[4] and Mean pooling (6.1%). On this dataset we also outperform Visual-Labels [27] (6.3%).

We report results on the **LSMDC** challenge[5], which combines M-VAD and MPII-MD. S2VT achieves 7.0% METEOR on the public test set using the evaluation server.

In Figure 4 we present descriptions generated by our model on some sample clips from the M-VAD dataset. More example video clips, generated sentences, and data are available on the authors' webpages[6].

6. Conclusion

This paper proposed a novel approach to video description. In contrast to related work, we construct descriptions using a sequence to sequence model, where frames are first read sequentially and then words are generated sequentially. This allows us to handle variable-length input and output while simultaneously modeling temporal structure. Our model achieves state-of-the-art performance on the MSVD dataset, and outperforms related work on two large and challenging movie-description datasets. Despite its conceptual simplicity, our model significantly benefits from additional data, suggesting that it has a high model capacity, and is able to learn complex temporal structure in the input and output sequences for challenging movie-description datasets.

Acknowledgments

Correct descriptions.

S2VT: A man is doing stunts on his bike.

S2VT: A herd of zebras are walking in a field.

S2VT: A young woman is doing her hair.

S2VT: A man is shooting a gun at a target.

(a)

Relevant but incorrect descriptions.

S2VT: A small bus is running into a building.

S2VT: A man is cutting a piece of a pair of a paper.

S2VT: A cat is trying to get a small board.

S2VT: A man is spreading butter on a tortilla.

(b)

Irrelevant descriptions.

S2VT: A man is pouring liquid in a pan.

S2VT: A polar bear is walking on a hill.

S2VT: A man is doing a pencil.
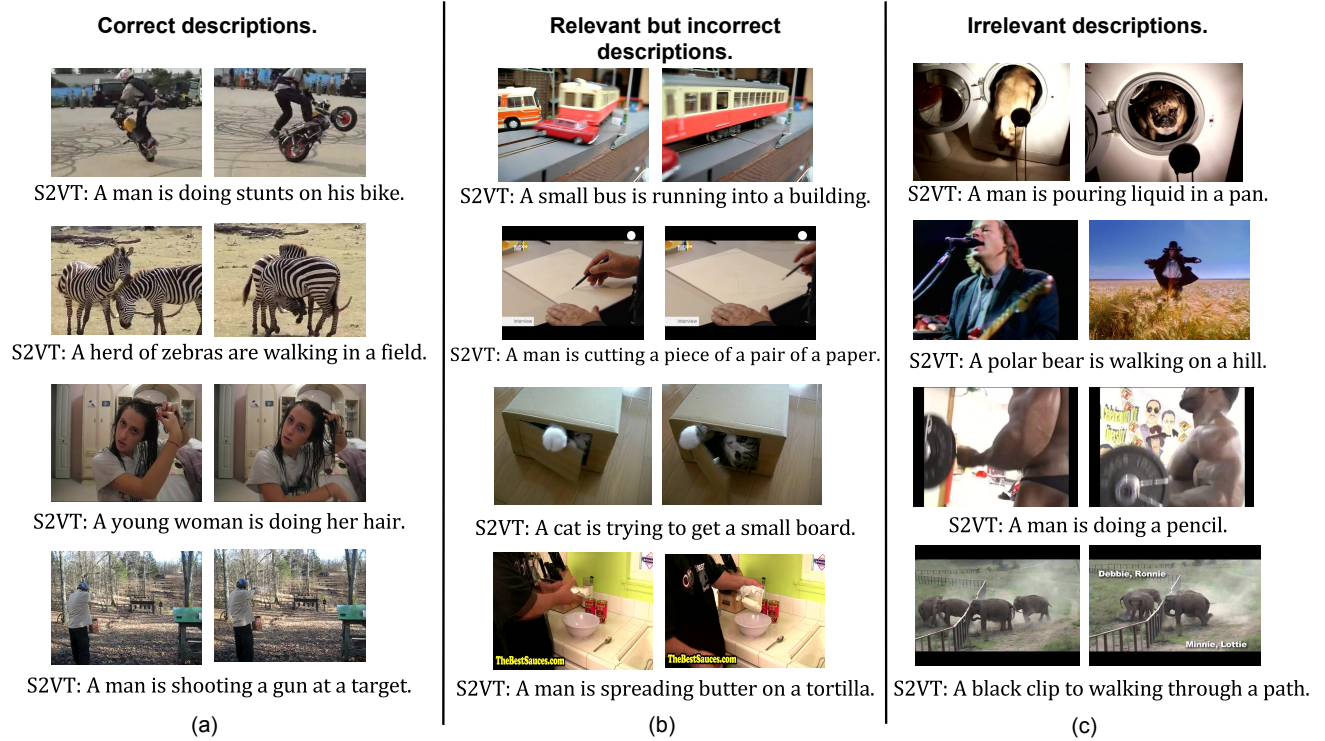
S2VT: A black clip to walking through a path.

(c)

Figure 3. Qualitative results on MSVD YouTube dataset from our S2VT model (RGB on VGG net). (a) Correct descriptions involving different objects and actions for several videos. (b) Relevant but incorrect descriptions. (c) Descriptions that are irrelevant to the event in the video.
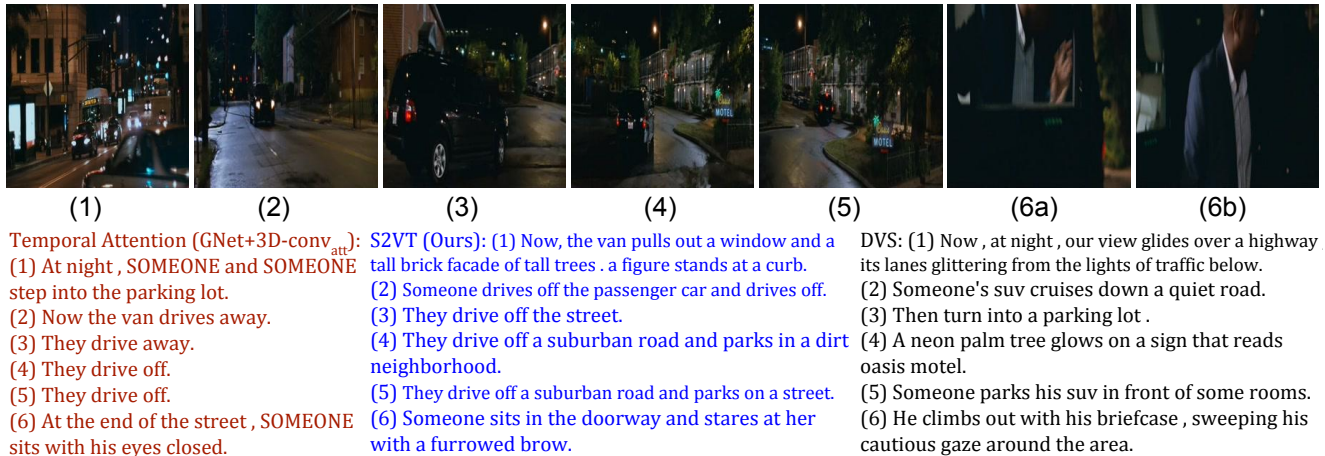

(1)   (2)   (3)   (4)   (5)   (6a)   (6b)

Temporal Attention (GNet+3D-conv_att):
(1) At night , SOMEONE and SOMEONE step into the parking lot.
(2) Now the van drives away.
(3) They drive away.
(4) They drive off.
(5) They drive off.
(6) At the end of the street , SOMEONE sits with his eyes closed.

S2VT (Ours): (1) Now, the van pulls out a window and a tall brick facade of tall trees . a figure stands at a curb.
(2) Someone drives off the passenger car and drives off.
(3) They drive off the street.
(4) They drive off a suburban road and parks in a dirt neighborhood.
(5) They drive off a suburban road and parks on a street.
(6) Someone sits in the doorway and stares at her with a furrowed brow.

DVS: (1) Now , at night , our view glides over a highway its lanes glittering from the lights of traffic below.
(2) Someone's suv cruises down a quiet road.
(3) Then turn into a parking lot .
(4) A neon palm tree glows on a sign that reads oasis motel.
(5) Someone parks his suv in front of some rooms.
(6) He climbs out with his briefcase , sweeping his cautious gaze around the area.

Figure 4. M-VAD Movie corpus: Representative frame from 6 contiguous clips from the movie "Big Mommas: Like Father, Like Son". From left: Temporal Attention (GoogleNet+3D-CNN) [43], S2VT (in blue) trained on the M-VAD dataset, and DVS: ground truth.

References

[1] H. Aradhye, G. Toderici, and J. Yagnik. Video2text: Learning to annotate video content. In *ICDMW*, 2009. 2

[2] T. Brox, A. Bruhn, N. Papenberg, and J. Weickert. High accuracy optical flow estimation based on a theory for warping. In *ECCV*, pages 25–36, 2004. 2, 4

[3] D. L. Chen and W. B. Dolan. Collecting highly parallel data for paraphrase evaluation. In *ACL*, 2011. 2, 5

[4] X. Chen, H. Fang, T.-Y. Lin, R. Vedantam, S. Gupta, P. Dol-

lar, and C. L. Zitnick. Microsoft COCO captions: Data collection and evaluation server. *arXiv:1504.00325*, 2015. 5

[5] X. Chen and C. L. Zitnick. Learning a recurrent visual representation for image caption generation. *CVPR*, 2015. 1

[6] K. Cho, B. van Merriënboer, D. Bahdanau, and Y. Bengio. On the properties of neural machine translation: Encoder-decoder approaches. *arXiv:1409.1259*, 2014. 3

[7] M. Denkowski and A. Lavie. Meteor universal: Language specific translation evaluation for any target language. In *EACL*, 2014. 5

[8] J. Donahue, L. A. Hendricks, S. Guadarrama, M. Rohrbach, S. Venugopalan, K. Saenko, and T. Darrell. Long-term recurrent convolutional networks for visual recognition and description. In *CVPR*, 2015. 1, 2, 3, 4

[9] G. Gkioxari and J. Malik. Finding action tubes. 2014. 4

[10] A. Graves and N. Jaitly. Towards end-to-end speech recognition with recurrent neural networks. In *ICML*, 2014. 1

[11] S. Guadarrama, N. Krishnamoorthy, G. Malkarnenkar, S. Venugopalan, R. Mooney, T. Darrell, and K. Saenko. Youtube2text: Recognizing and describing arbitrary activities using semantic hierarchies and zero-shoot recognition. In *ICCV*, 2013. 1, 2

[12] S. Hochreiter and J. Schmidhuber. Long short-term memory. *Neural computation*, 9(8), 1997. 1, 3

[13] P. Hodosh, A. Young, M. Lai, and J. Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. In *TACL*, 2014. 6

[14] H. Huang, Y. Lu, F. Zhang, and S. Sun. A multi-modal clustering method for web videos. In *ISCTCS*. 2013. 2

[15] Y. Jia, E. Shelhamer, J. Donahue, S. Karayev, J. Long, R. Girshick, S. Guadarrama, and T. Darrell. Caffe: Convolutional architecture for fast feature embedding. *ACMMM*, 2014. 2

[16] A. Karpathy and L. Fei-Fei. Deep visual-semantic alignments for generating image descriptions. *CVPR*, 2015. 1

[17] D. Kingma and J. Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014. 7

[18] R. Kiros, R. Salakhutdinov, and R. S. Zemel. Unifying visual-semantic embeddings with multimodal neural language models. *arXiv:1411.2539*, 2014. 1

[19] N. Krishnamoorthy, G. Malkarnenkar, R. J. Mooney, K. Saenko, and S. Guadarrama. Generating natural-language video descriptions using text-mined knowledge. In *AAAI*, July 2013. 2

[20] P. Kuznetsova, V. Ordonez, T. L. Berg, U. C. Hill, and Y. Choi. Treetalk: Composition and compression of trees for image descriptions. In *TACL*, 2014. 1

[21] C.-Y. Lin. Rouge: A package for automatic evaluation of summaries. In *Text Summarization Branches Out: Proceedings of the ACL-04 Workshop*, pages 74–81, 2004. 5

[22] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick. Microsoft coco: Common objects in context. In *ECCV*, 2014. 6

[23] J. Mao, W. Xu, Y. Yang, J. Wang, and A. L. Yuille. Deep captioning with multimodal recurrent neural networks (m-rnn). *arXiv:1412.6632*, 2014. 1

[24] J. Y. Ng, M. J. Hausknecht, S. Vijayanarasimhan, O. Vinyals, R. Monga, and G. Toderici. Beyond short snippets: Deep networks for video classification. *CVPR*, 2015. 3, 4

[25] P. Over, G. Awad, M. Michel, J. Fiscus, G. Sanders, B. Shaw, A. F. Smeaton, and G. Quéenot. TRECVID 2012 – an overview of the goals, tasks, data, evaluation mechanisms and metrics. In *Proceedings of TRECVID 2012*, 2012. 2

[26] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu. Bleu: a method for automatic evaluation of machine translation. In *ACL*, 2002. 5

[27] A. Rohrbach, M. Rohrbach, and B. Schiele. The long-short story of movie description. *GCPR*, 2015. 7

[28] A. Rohrbach, M. Rohrbach, N. Tandon, and B. Schiele. A dataset for movie description. In *CVPR*, 2015. 1, 2, 5, 7

[29] M. Rohrbach, W. Qiu, I. Titov, S. Thater, M. Pinkal, and B. Schiele. Translating video content to natural language descriptions. In *ICCV*, 2013. 1

[30] O. Russakovsky, J. Deng, H. Su, J. Krause, S. Satheesh, S. Ma, Z. Huang, A. Karpathy, A. Khosla, M. Bernstein, A. C. Berg, and L. Fei-Fei. ILSVRC, 2014. 4

[31] K. Simonyan and A. Zisserman. Two-stream convolutional networks for action recognition in videos. In *NIPS*, 2014. 2

[32] K. Simonyan and A. Zisserman. Very deep convolutional networks for large-scale image recognition. *CoRR*, abs/1409.1556, 2014. 4

[33] N. Srivastava, E. Mansimov, and R. Salakhutdinov. Unsupervised learning of video representations using LSTMs. *ICML*, 2015. 2

[34] I. Sutskever, O. Vinyals, and Q. V. Le. Sequence to sequence learning with neural networks. In *NIPS*, 2014. 1, 2, 3

[35] C. Szegedy, W. Liu, Y. Jia, P. Sermanet, S. Reed, D. Anguelov, D. Erhan, V. Vanhoucke, and A. Rabinovich. Going deeper with convolutions. *CVPR*, 2015. 6

[36] J. Thomason, S. Venugopalan, S. Guadarrama, K. Saenko, and R. J. Mooney. Integrating language and vision to generate natural language descriptions of videos in the wild. In *COLING*, 2014. 2, 6

[37] A. Torabi, C. Pal, H. Larochelle, and A. Courville. Using descriptive video services to create a large data source for video annotation research. *arXiv:1503.01070v1*, 2015. 2, 5

[38] R. Vedantam, C. L. Zitnick, and D. Parikh. CIDEr: Consensus-based image description evaluation. *CVPR*, 2015. 5

[39] S. Venugopalan, H. Xu, J. Donahue, M. Rohrbach, R. Mooney, and K. Saenko. Translating videos to natural language using deep recurrent neural networks. In *NAACL*, 2015. 1, 2, 4, 5, 6, 7

[40] O. Vinyals, A. Toshev, S. Bengio, and D. Erhan. Show and tell: A neural image caption generator. *CVPR*, 2015. 1, 2, 4

[41] H. Wang and C. Schmid. Action recognition with improved trajectories. In *ICCV*, pages 3551–3558. IEEE, 2013. 2

[42] S. Wei, Y. Zhao, Z. Zhu, and N. Liu. Multimodal fusion for video search reranking. *TKDE*, 2010. 2

[43] L. Yao, A. Torabi, K. Cho, N. Ballas, C. Pal, H. Larochelle, and A. Courville. Describing videos by exploiting temporal structure. *arXiv:1502.08029v4*, 2015. 1, 2, 4, 6, 7, 8

[44] W. Zaremba and I. Sutskever. Learning to execute. *arXiv:1410.4615*, 2014. 3